

CS304 Operating Systems

DR GAYATHRI ANANTHANARAYANAN

gayathri@iitdh.ac.in

Materials in these slides have been borrowed from textbooks and existing operating systems courses

13th Jan 2021, Mechanism of Process Execution

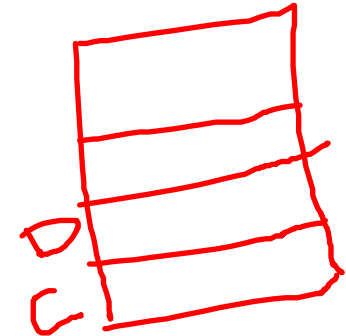
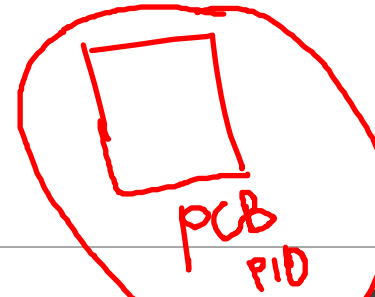
How does the OS run a process?

How does it handle a system call?

How does it context switch from one process to the other?

Process Execution

- OS creates a process entry for the process in a process list
- Allocates memory and creates memory image
 - Code and data (from executable)
 - Stack and heap
- Points CPU program counter to current instruction
 - Other registers may store operands, return values etc.
- After setup, OS is out of the way and process executes directly on CPU



A simple function call

User



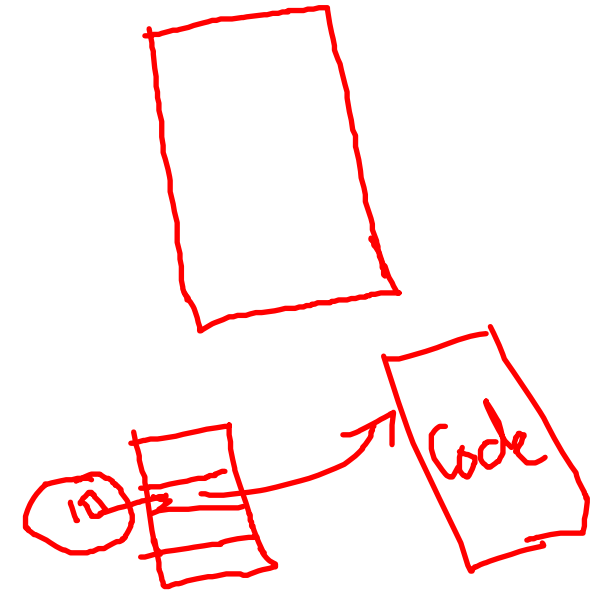
- A function call translates to a jump instruction
- A new stack frame pushed to stack and stack pointer (SP) updated
- Old value of PC (return value) pushed to stack and PC updated
- Stack frame contains return value, function arguments etc.

How is a system call different?

CPU hardware has multiple privilege levels

- One to run user code: user mode
- One to run OS code like system calls: kernel mode
- Some instructions execute only in kernel mode
- Kernel does not trust user stack
- Uses a separate kernel stack when in kernel mode
- Kernel does not trust user provided addresses to jump to
- Kernel sets up Interrupt Descriptor Table (IDT) at boot time– IDT has addresses of kernel functions to run for system calls and other events

open()
(Josef)
↓



Mechanism of system call: trap instruction

trap

ISA

↳ x86
x86-64

When system call must be made, a special trap instruction is run (usually hidden from user by libc)

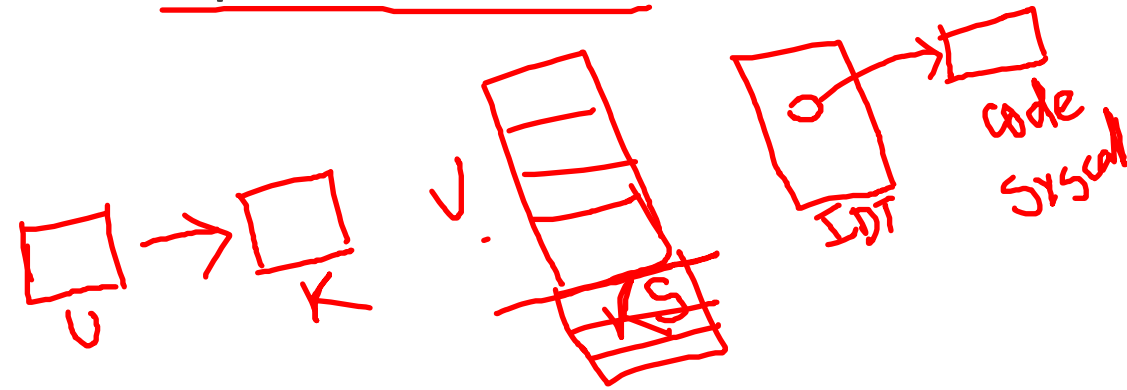
- Trap instruction execution

- Move CPU to higher privilege level

- Switch to kernel stack

- Save context (old PC, registers) on kernel stack

- Look up address in IDT and jump to trap handler function in OS code



More on trap instruction

Trap instruction is executed on hardware in following cases:

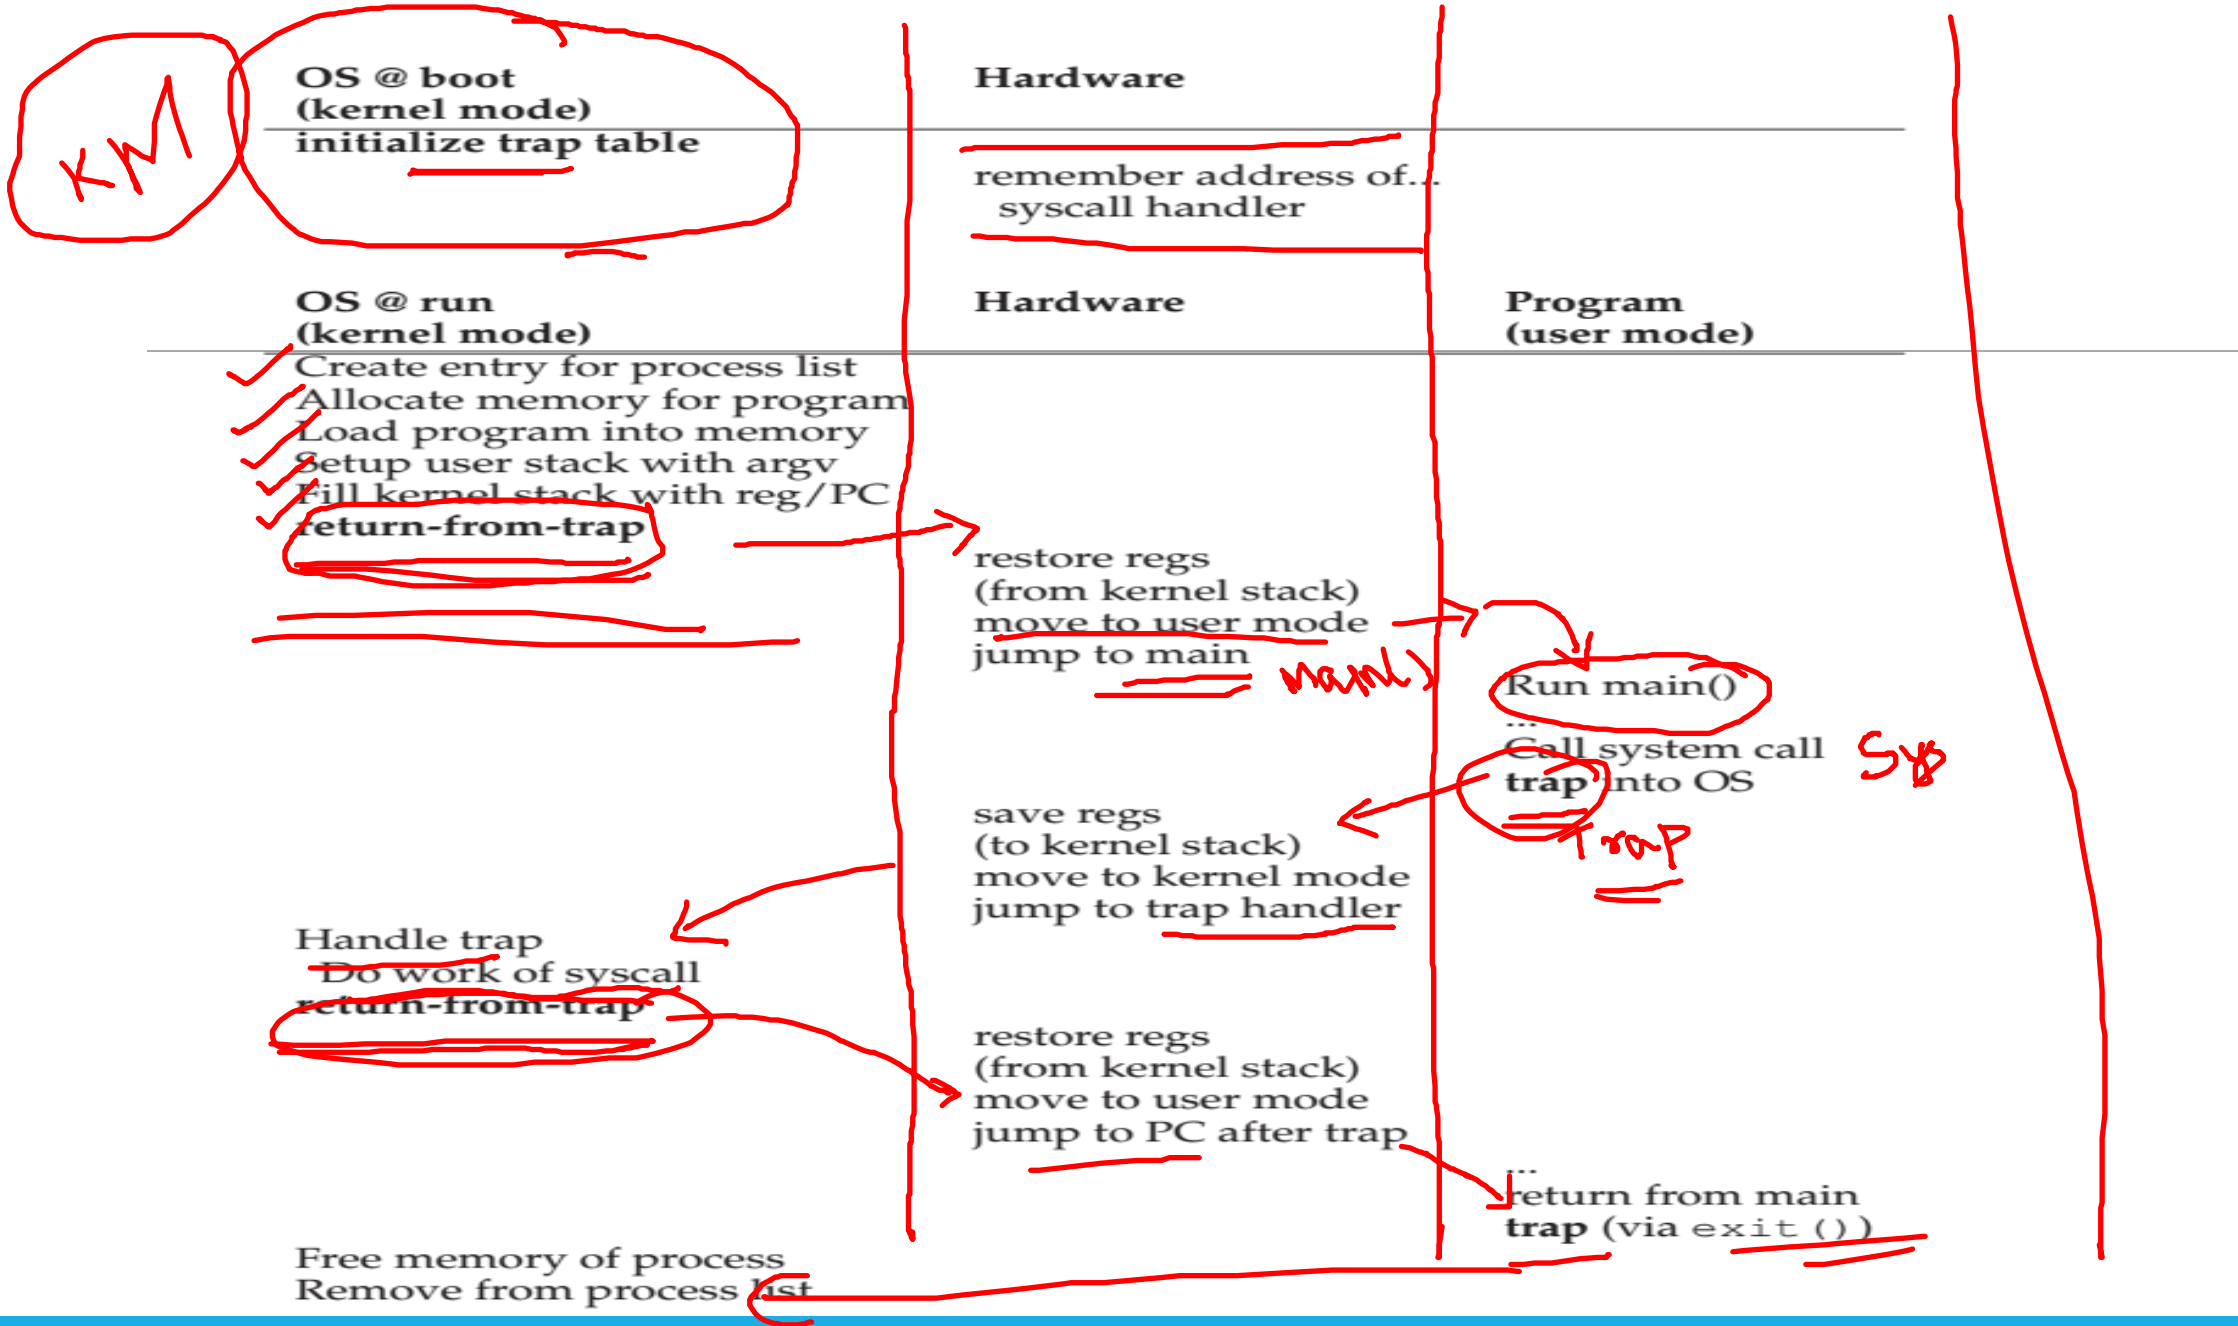
- System call (program needs OS service)
- Program fault (program does something illegal, e.g., access memory it doesn't have access to)
- Interrupt (external device needs attention of OS, e.g., a network packet has arrived on network card)
- Across all cases, the mechanism is: save context on kernel stack and switch to OS address in IDT
- IDT has many entries: which to use?
 - System calls/interrupts store a number in a CPU register before calling trap, to identify which IDT entry to use

Return from trap

When OS is done handling syscall or interrupt, it calls a special instruction return-from-trap

- Restore context of CPU registers from kernel stack
- Change CPU privilege from kernel mode to user mode
- Restore PC and jump to user code after trap
- User process unaware that it was suspended, resumes execution as always
- Must you always return to the same user process from kernel mode? No
- Before returning to user mode, OS checks if it must switch to another process

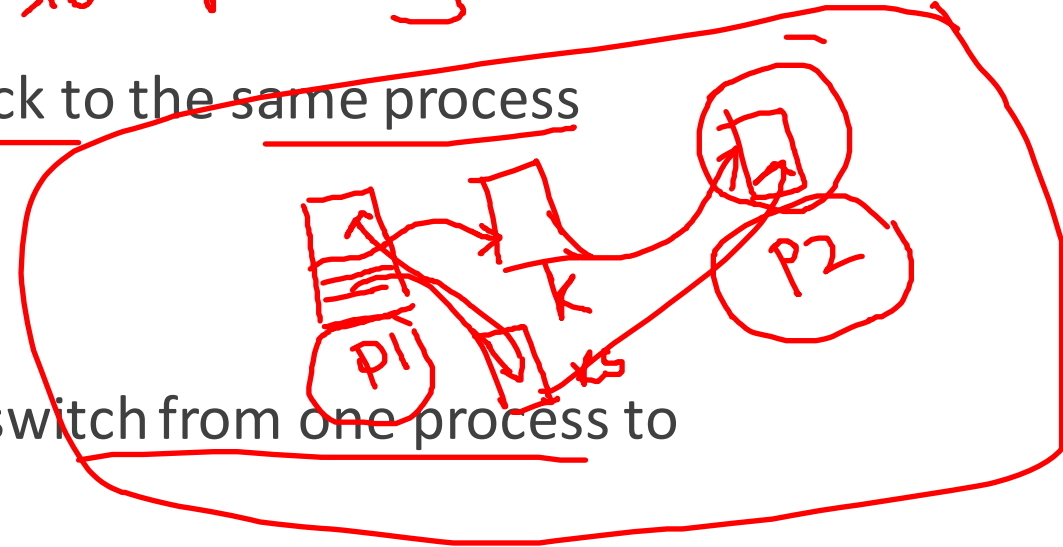




Why switch between processes?

- Sometimes when OS is in kernel mode, it cannot return back to the same process it left
 - Process has exited or must be terminated (e.g., segfault)
 - Process has made a blocking system call
- Sometimes, the OS does not want to return back to the same process
 - The process has run for too long ✓
 - Must timeshare CPU with other processes
- In such cases, OS performs a context switch to switch from one process to another

I/O Synchronous?



The OS scheduler

- OS scheduler has two parts
 - Policy to pick which process to run (next lecture)
 - Mechanism to switch to that process (this lecture)
- Non preemptive (cooperative) schedulers are polite
 - Switch only if process blocked or terminated
- Preemptive (non-cooperative) schedulers can switch even when process is ready to continue
 - CPU generates periodic timer interrupt
 - After servicing interrupt, OS checks if the current process has run for too long

*infinite while(1) {
loop }*

XMS

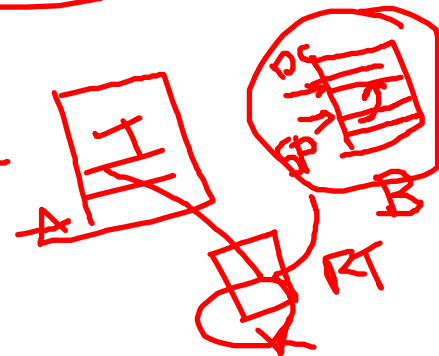
B

A

Context Switch Mechanism

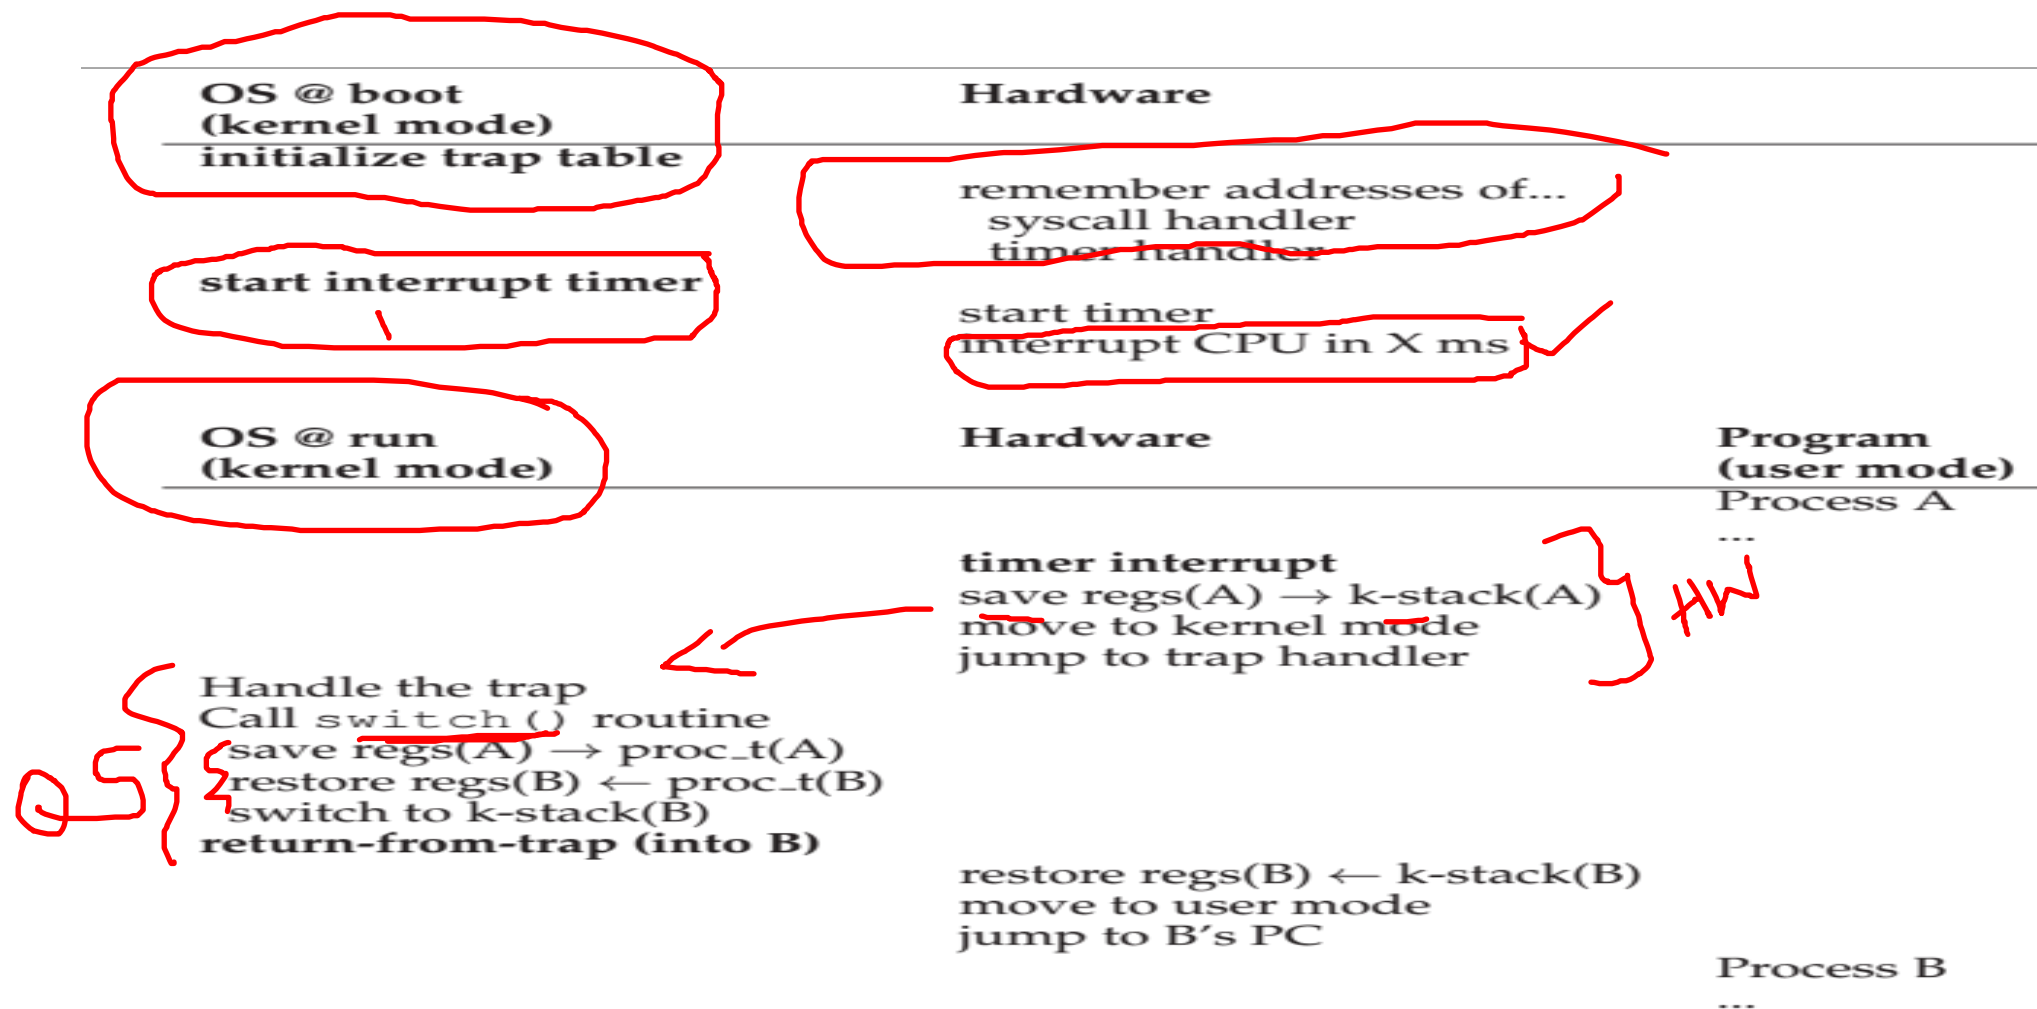
Example: process A has moved from user to kernel mode, OS decides it must switch from A to B

- Save context (PC, registers, kernel stack pointer) of A on kernel stack
- Switch SP to kernel stack of B
- Restore context from B's kernel stack
- Who has saved registers on B's kernel stack? — OS did, when it switched out B in the past
- Now, CPU is running B in kernel mode, return-from-trap to switch to user mode of B



A subtlety on saving context

- Context (PC and other CPU registers) saved on the kernel stack in two different scenarios
- When going from user mode to kernel mode, user context (e.g., which instruction of user code you stopped at) is saved on kernel stack by the trap instruction
 - Restored by return-from-trap
- During a context switch, kernel context (e.g., where you stopped in the OS code) of process A is saved on the kernel stack of A by the contextswitching code
 - Restores kernel context of process B



Trivia

What is the kernel stack of a process used for? ✓

When a process goes from user mode to kernel mode due to a trap occurring (system call / interrupt / program fault), who saves the context of the process? What exactly happens? ✓

What happens on a trap? ✓

Who saves user's CPU context when moving from user mode to kernel mode? ✓

Are different types of contexts saved on the kernel stack? ✓

```
# void swtch(struct context **old, struct context *new);
#
# Save current register context in old
# and then load register context from new.
.globl swtch
swtch:
    # Save old registers
    movl 4(%esp), %eax # put old ptr into eax
    popl 0(%eax)       # save the old IP
    movl %esp, 4(%eax) # and stack
    movl %ebx, 8(%eax) # and other registers
    movl %ecx, 12(%eax)
    movl %edx, 16(%eax)
    movl %esi, 20(%eax)
    movl %edi, 24(%eax)
    movl %ebp, 28(%eax)

    # Load new registers
    movl 4(%esp), %eax # put new ptr into eax
    movl 28(%eax), %ebp # restore other registers
    movl 24(%eax), %edi
    movl 20(%eax), %esi
    movl 16(%eax), %edx
    movl 12(%eax), %ecx
    movl 8(%eax), %ebx
    movl 4(%eax), %esp # stack is switched here
    pushl 0(%eax)      # return addr put in place
    ret               # finally return into new ctxt
```

We thus have the basic mechanisms for virtualizing the CPU in place.

But a major question is left unanswered: which process should we run at a given time?

It is this question that the scheduler must answer, and thus the next topic of our study.