

CS304 Operating Systems

DR GAYATHRI ANANTHANARAYANAN

gayathri@iitdh.ac.in

Materials in these slides have been borrowed from textbooks and existing operating systems courses

21 Jan 2021 Inter Process Communication

- Processes do not share any memory with each other
- Some processes might want to work together for a task, so need to communicate information
- IPC mechanisms to share information between processes

Virtual Memory View

- During execution, each process can only view its virtual addresses,
- It cannot
 - View another processes virtual address space
 - Determine the physical address mapping

Executing
Process

Virtual Memory Map

6
5
4
3
2
1

Virtual Memory Map

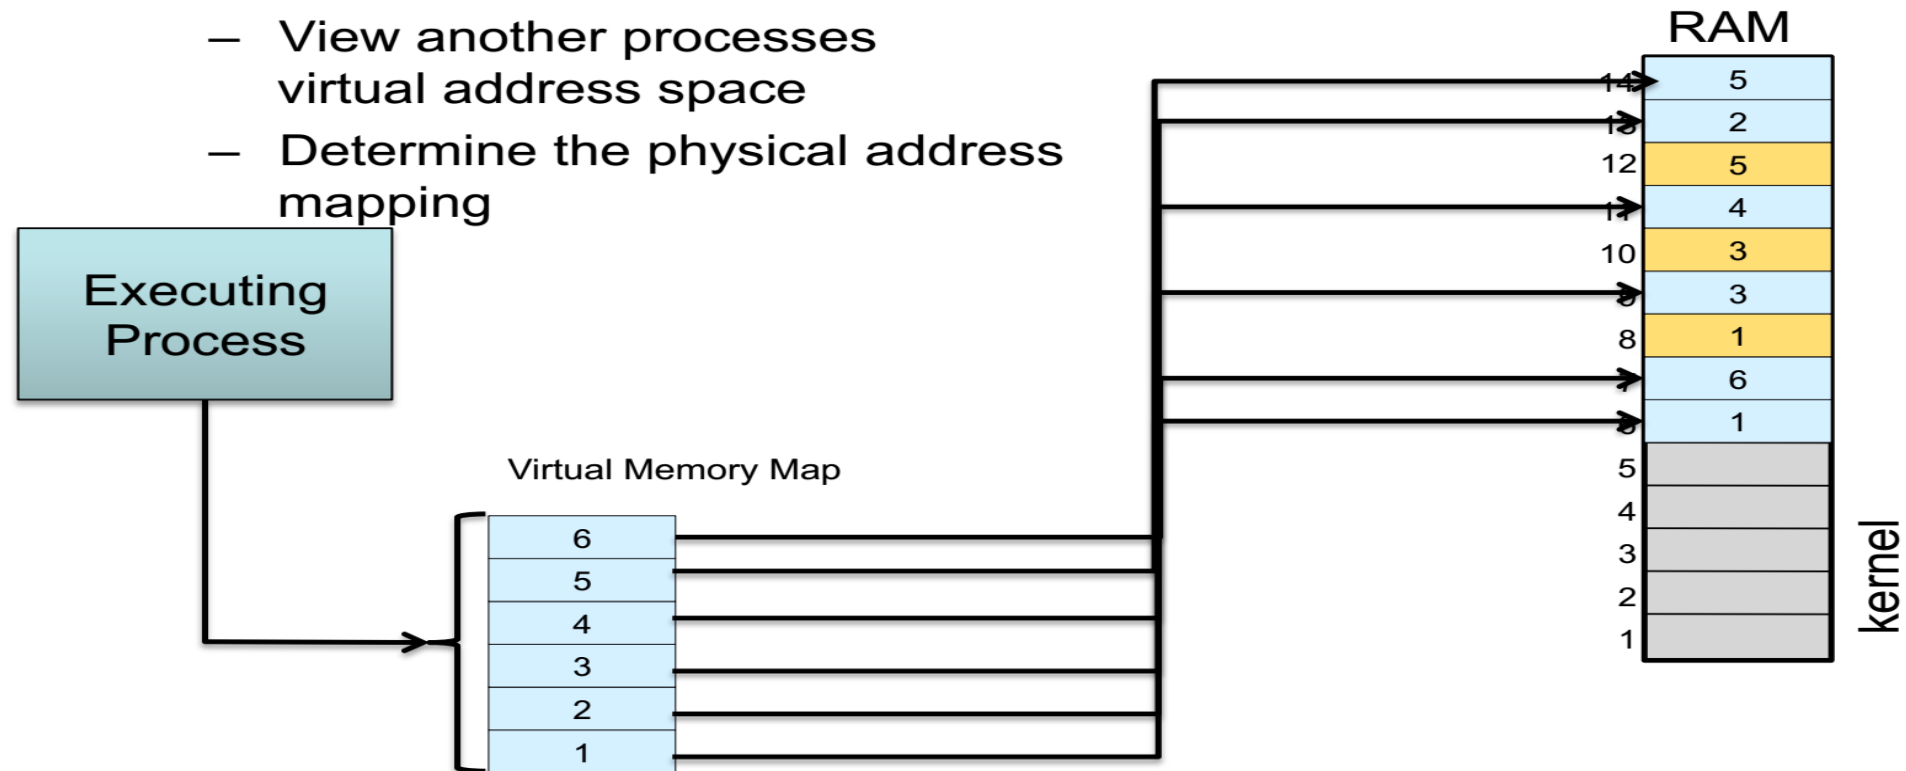
6
5
4
3
2
1

RAM

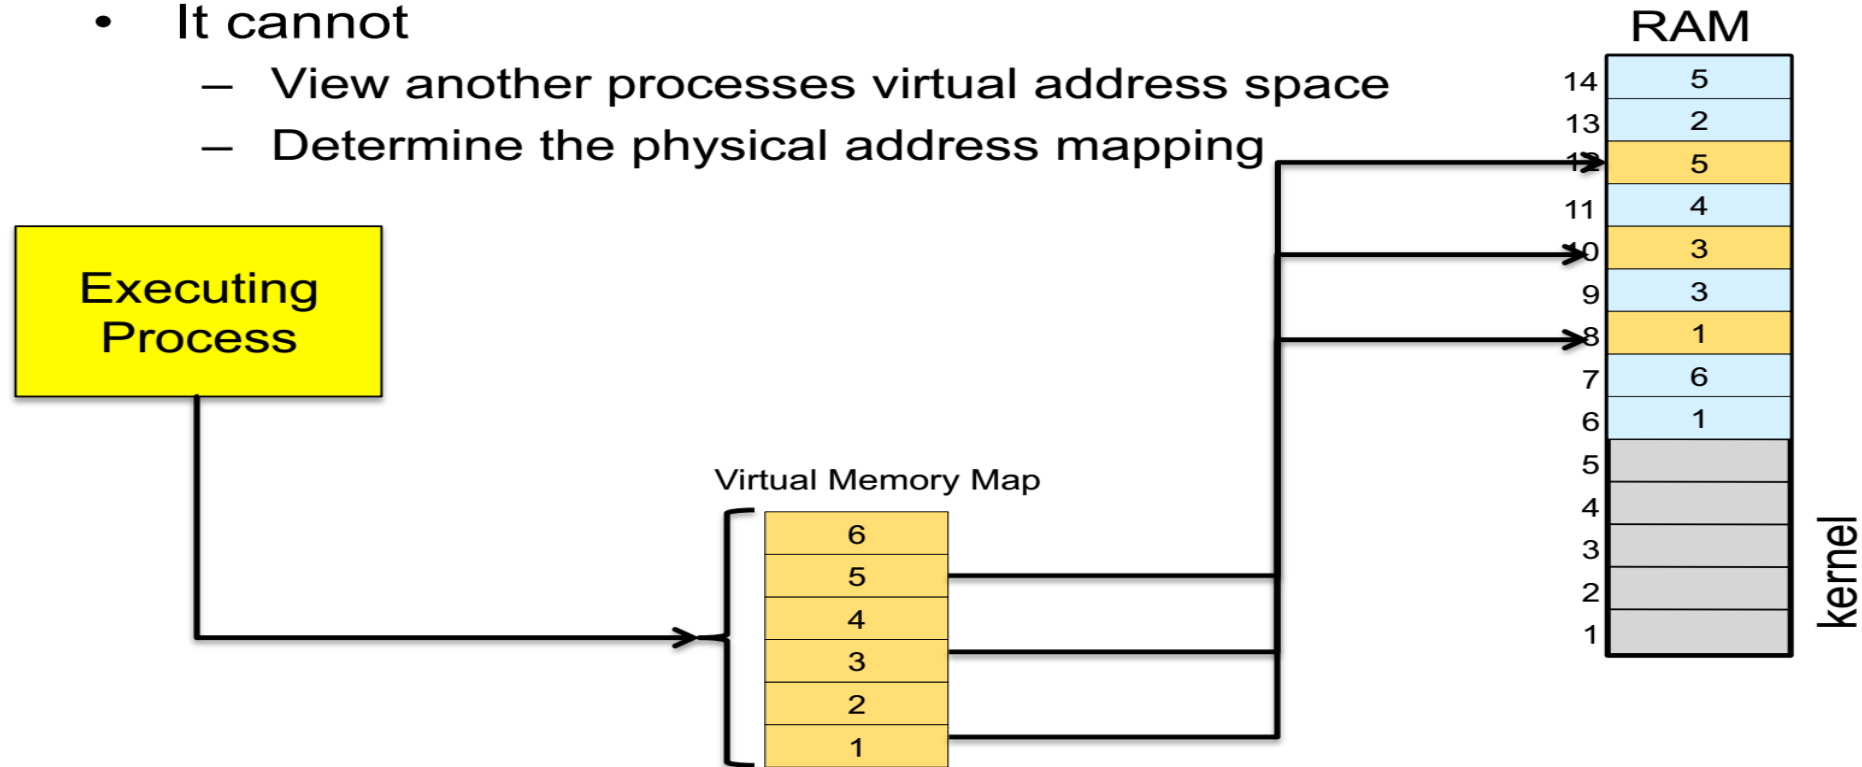
14	5
13	2
12	5
11	4
10	3
9	3
8	1
7	6
6	1
5	
4	
3	
2	
1	

kernel

- During execution, each process can only view its virtual addresses
- It cannot
 - View another processes virtual address space
 - Determine the physical address mapping



- During execution, each process can only view its virtual addresses,
- It cannot
 - View another processes virtual address space
 - Determine the physical address mapping




Inter Process Communication

- Advantages of Inter Process Communication (IPC)
 - Information sharing
 - Modularity/Convenience
- 3 ways
 - Shared memory
 - Message Passing
 - Signals

Shared Memory

- One process will create an area in RAM which the other process can access
- Both processes can access shared memory like a regular working memory
 - Reading/writing is like regular reading/writing
 - Fast
- **Limitation** : Error prone. Needs synchronization between processes



Shared Memory

Processes can both access same region of memory via `shmget()` system call

- `int shmget(key_t key, int size, int shmflg)`
- By providing same key, two processes can get same segment of memory
- Can read/write to memory to communicate
- Need to take care that one is not overwriting other's data: how?

IPC - Shared Memory
IPC - Shared Memory

Shared Memory in Linux

- **int shmget (key, size, flags)**
 - Create a shared memory segment;
 - Returns ID of segment : **shmid**
 - **key** : unique identifier of the shared memory segment
 - **size** : size of the shared memory (rounded up to the PAGE_SIZE)
- **int shmat(shmid, addr, flags)**
 - **A**ttach **shmid** shared memory to address space of the calling process
 - **addr** : pointer to the shared memory address space ✓
- **int shmdt(shmid)**
 - **Det**ach shared memory

Example

server.c

```
1 #include <sys/types.h>
2 #include <sys/ipc.h>
3 #include <sys/shm.h>
4 #include <stdio.h>
5 #include <stdlib.h>
6
7 #define SHMSIZE 27 /* Size of shared memory */
8
9 main()
10 {
11     char c;
12     int shmid;
13     key_t key;
14     char *shm, *s;
15
16     key = 5678; /* some key to uniquely identifies the shared memory */
17
18     /* Create the segment. */
19     if ((shmid = shmget(key, SHMSIZE, IPC_CREAT | 0666)) < 0) {
20         perror("shmget");
21         exit(1);
22     }
23
24     /* Attach the segment to our data space. */
25     if ((shm = shmat(shmid, NULL, 0)) == (char *) -1) {
26         perror("shmat");
27         exit(1);
28     }
29
30     /* Now put some things into the shared memory */
31     s = shm;
32     for (c = 'a'; c <= 'z'; c++)
33         *s++ = c;
34     *s = 0; /* end with a NULL termination */
35
36     /* Wait until the other process changes the first character
37      * to '*' the shared memory */
38     while (*shm != '*')
39         sleep(1);
40     exit(0);
41 }
```

client.c

```
1 #include <sys/types.h>
2 #include <sys/ipc.h>
3 #include <sys/shm.h>
4 #include <stdio.h>
5 #include <stdlib.h>
6
7 #define SHMSIZE 27
8
9 main()
10 {
11     int shmid;
12     key_t key;
13     char *shm, *s;
14
15     /* We need to get the segment named "5678", created by the server
16     key = 5678;
17
18     /* Locate the segment. */
19     if ((shmid = shmget(key, SHMSIZE, 0666)) < 0) {
20         perror("shmget");
21         exit(1);
22     }
23
24     /* Attach the segment to our data space. */
25     if ((shm = shmat(shmid, NULL, 0)) == (char *) -1) {
26         perror("shmat");
27         exit(1);
28     }
29
30     /* read what the server put in the memory. */
31     for (s = shm; *s != 0; s++)
32         putchar(*s);
33     putchar('\n');
34
35     /*
36     * Finally, change the first character of the
37     * segment to '*', indicating we have read
38     * the segment.
39     */
40     *shm = '*';
41
42     exit(0);
43 }
```

Signals

- A certain set of signals supported by OS
 - Some signals have fixed meaning (e.g., signal to terminate process)
 - Some signals can be user-defined
- Signals can be sent to a process by OS or another process (e.g., if you type Ctrl+C, OS sends SIGINT signal to running process)
- Signal handler: every process has a default code to execute for each signal
 - Exit on terminate signal
- Some signal handlers can be overridden to do other things

-
- Asynchronous unidirectional communication between processes
 - Signals are a small integer
 - eg. 9: kill, 11: segmentation fault
 - Send a signal to a process
 - `kill(pid, signum)` ✓
 - Process handler for a signal
 - `sighandler_t signal(signum, handler);`
 - Default if no handler defined

Sockets



- Sockets can be used for two processes on same machine or different machines to communicate
 - TCP/UDP sockets across machines
 - Unix sockets in local machine
- Communicating with sockets
 - Processes open sockets and connect them to each other
 - Messages written into one socket can be read from another
 - OS transfers data across socket buffers

Pipes

- Pipe system call returns two file descriptors
 - Read handle and write handle
 - A pipe is a half-duplex communication
 - Data written in one file descriptor can be read through another
- Regular pipes: both fd are in same process (how it is useful?)
 - Parent and child share fd after fork
 - Parent uses one end and child uses other end
- Named pipes: two endpoints of a pipe can be in different processes
- Pipe data buffered in OS buffers between write and read



fork()

Half duplex
pipe()




Pipes


- Always between parent and child
- Always unidirectional
- Accessed by two associated file descriptors:
 - fd[0] for reading from pipe
 - fd[1] for writing to the pipe



Two Way Communication



- Two pipes opened
pipe0 and pipe1
- Note the unnecessary
pipes



- Close the unnecessary
pipes

Example [Child process sending a string to Parent]

```
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>

int main(){
    int pipefd[2];
    int pid;
    char recv[32];

    pipe(pipefd);

    switch(pid=fork()) {
    case -1: perror("fork");
            exit(1);
    case 0:
        close(pipefd[0]);
        FILE *out = fdopen(pipefd[1], "w");
        fprintf(out, "Hello World\n");
        break;
    default:
        close(pipefd[1]);
        FILE *in = fdopen(pipefd[0], "r");
        fscanf(in, "%s", recv);
        printf("%s", recv);
        break;
    }
}
```

dup(1)
dup(2)

Named
FIFOs

Message Queues

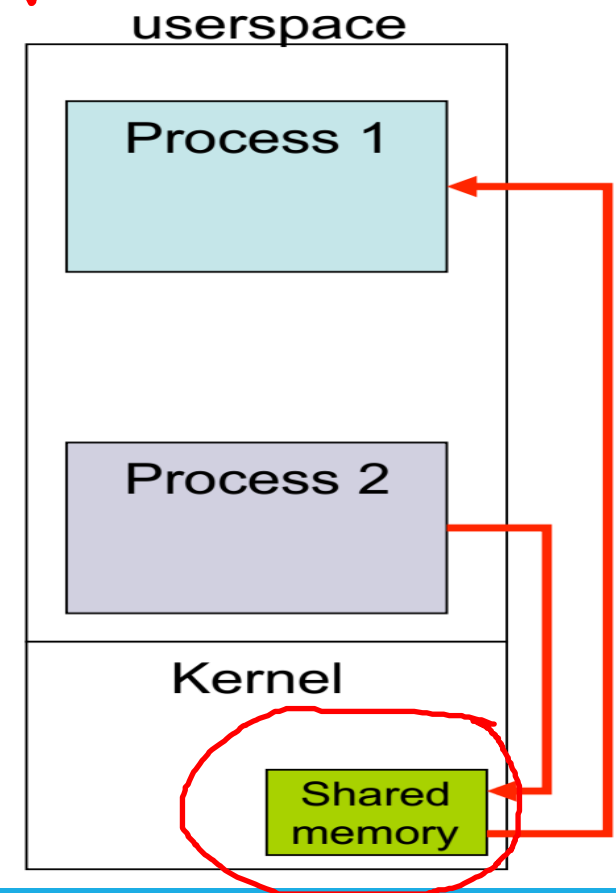
- Mailbox abstraction
- Process can open a mailbox at a specified location
- Processes can send/receive messages from mailbox
- OS buffers messages between send and receive

MB

Message Passing ✓

*msgget ()
msgsend ()*

- Shared memory created in the kernel
- System calls such as send and receive used for communication
 - Cooperating : each send must have a receive
- **Advantage** : Explicit sharing, less error prone
- **Limitation** : Slow. Each call involves marshalling / demarshalling of information



Blocking Vs Non-Blocking

- Some IPC actions can block
 - Reading from socket/pipe that has no data, or reading from empty message queue
 - Writing to a full socket/pipe/message queue
- The system calls to read/write have versions that block or can return with an error code in case of failure
 - A socket read can return error indicating no data to be read, instead of blocking