

CS304 Operating Systems

DR GAYATHRI ANANTHANARAYANAN

gayathri@iitdh.ac.in

Materials in these slides have been borrowed from textbooks and existing operating systems courses

A solid blue horizontal bar spanning the width of the slide, located at the bottom.


Jan27, 2021 – Threads & Concurrency

So, far we have studied single threaded programs

- Recall: Process execution
 - PC points to current instruction being run
 - SP points to stack frame of current function call
- A program can also have multiple threads of execution
- What is a thread?

Process Execution

- Separate streams of execution
- Each process isolated from the other
- Process state contains
 - Process ID ✓
 - Environment ✓
 - Working directory. ✓
 - Program instructions ✓
 - Registers ✓
 - Stack ✓
 - Heap ✓
 - File descriptors ✓
- Created by the OS using fork
 - Significant overheads



Threads


A thread is like another copy of a process that executes independently

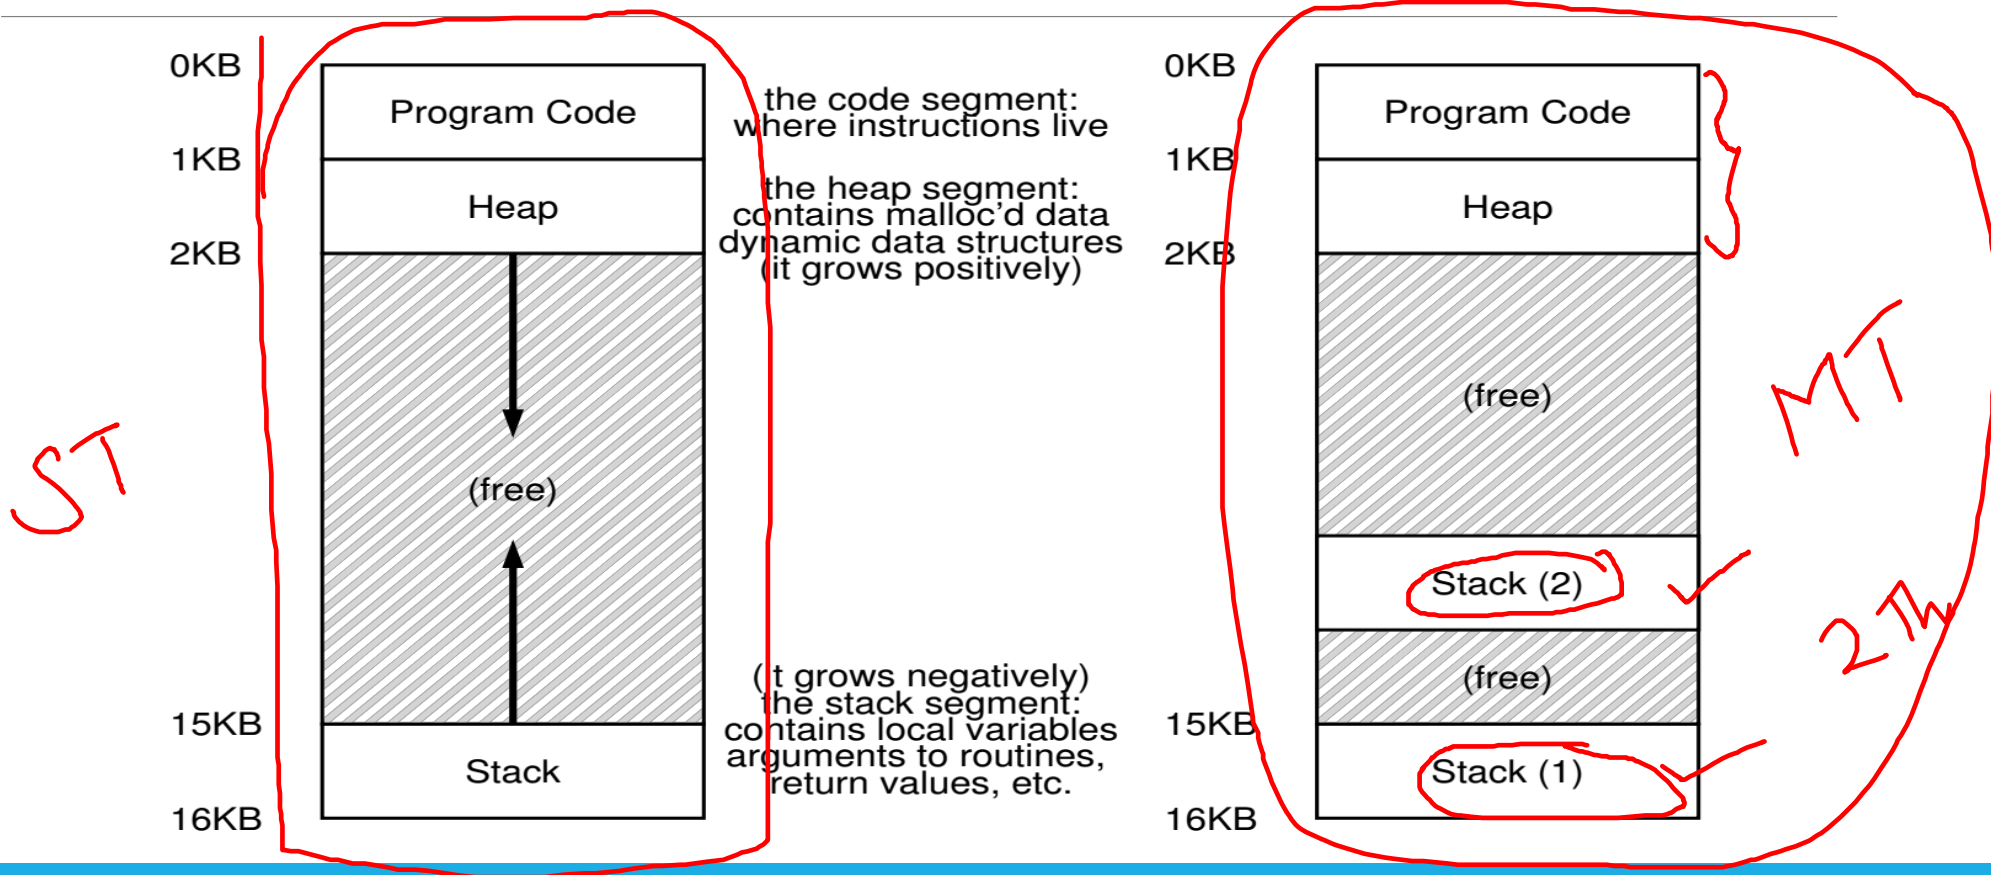
- Threads shares the same address space (code, heap)
- Each thread has separate PC ← *Private Regs, Stack*
- Each thread may run over different part of the program
- Each thread has separate stack for independent function calls

Threads


- Separate streams of execution within a single process
- Threads in a process not isolated from each other
- Each thread state (thread control block) contains
 - Registers (including EIP, ESP)
 - stack

TCLB





Why threads ?



- Parallelism: a single process can effectively utilize multiple CPU cores
- Understand the difference between concurrency and parallelism
- Concurrency: running multiple threads/processes at the same time, even on single CPU core, by interleaving their executions
- Parallelism: running multiple threads/processes in parallel over different CPU cores
- Even if no parallelism, concurrency of threads ensures effective use of CPU when one of the threads blocks (e.g., for I/O)

TRUE

Why threads?

- Lightweight

Platform	fork()			pthread_create()		
	real	user	sys	real	user	sys
Intel 2.6 GHz Xeon E5-2670 (16 cores/node)	8.1	0.1	2.9	0.9	0.2	0.3
Intel 2.8 GHz Xeon 5660 (12 cores/node)	4.4	0.4	4.3	0.7	0.2	0.5
AMD 2.3 GHz Opteron (16 cores/node)	12.5	1.0	12.5	1.2	0.2	1.3
AMD 2.4 GHz Opteron (8 cores/node)	17.6	2.2	15.7	1.4	0.3	1.3
IBM 4.0 GHz POWER6 (8 cpus/node)	9.5	0.6	8.8	1.6	0.1	0.4
IBM 1.9 GHz POWER5 p5-375 (8 cpus/node)	64.2	30.7	27.6	1.7	0.6	1.1
IBM 1.5 GHz POWER4 (8 cpus/node)	104.5	48.6	47.2	2.1	1.0	1.5
INTEL 2.4 GHz Xeon (2 cpus/node)	54.9	1.5	20.8	1.6	0.7	0.9
INTEL 1.4 GHz Itanium2 (4 cpus/node)	54.5	1.1	22.2	2.0	1.2	0.6

Cost of creating 50,000 processes / threads

(<https://computing.llnl.gov/tutorials/pthreads/>)

- Efficient communication between entities
- Efficient context switching

Threads vs Processes

Parent P forks a child C

- P and C do not share any memory
- Need complicated IPC mechanisms to communicate
- Extra copies of code, data in memory
- Parent P executes two threads T1 and T2
 - T1 and T2 share parts of the address space
 - Global variables can be used for communication
 - Smaller memory footprint
- Threads are like separate processes, except they share the same address space

Amdahl's Law

Comparison

- | | |
|---|--|
| <ul style="list-style-type: none">• A thread has no data segment or heap• A thread cannot live on its own. It needs to be attached to a process• There can be more than one thread in a process. Each thread has its own stack• If a thread dies, its stack is reclaimed | <ul style="list-style-type: none">• A process has code, heap, stack, other segments• A process has at-least one thread.• Threads within a process share the same I/O, code, files.• If a process dies, all threads die. |
|---|--|

Trivia

Do the following statements apply to [✓]Processes (P), [✓]Threads (T) or both (B)

Can share a virtual address space

— Threads

Take longer to context switch

— P

Have an execution context


— Both

Usually result in hotter caches when multiple exist

Make use of some communication mechanisms

— B — T

Scheduling Threads



- OS schedules threads that are ready to run independently, much like processes
- The context of a thread (PC, registers) is saved into/restored from thread control block (TCB) ✓
 - Every PCB has one or more linked TCBs
- Threads that are scheduled independently by kernel are called kernel threads
- In contrast, some libraries provide user-level threads
 - User program sees multiple threads ✓
 - Library multiplexes larger number of user threads over a smaller number of kernel threads
 - Low overhead of switching between user threads (no expensive context switch)
 - But multiple user threads cannot run in parallel -2

posix ✗

pthread library

- Create a thread in a process

int pthread_create(pthread_t *thread, ✓

const pthread_attr_t *attr,

void *(*start_routine) (void *),

void *arg);

Thread identifier (TID) much like

Pointer to a function,
which starts execution in a
different thread

Arguments to the function

- ~~Destroying a thread~~

void pthread_exit(void *retval);

Exit value of the thread

Pthread contd..

pthread_detach()

- Join : Wait for a specific thread to complete

int pthread_join(pthread_t thread, void **retval);

TID of the thread to wait for

Exit status of the thread

what is the difference with wait()?

Example

```

1 #include <stdio.h>
2 #include <assert.h>
3 #include <pthread.h>
4 #include "common.h"
5 #include "common_threads.h"
6
7 void *mythread(void *arg) {
8     printf("%s\n", (char *) arg);
9     return NULL;
10 }
11
12 int
13 main(int argc, char *argv[]) {
14     pthread_t p1, p2;
15     int rc;
16     printf("main: begin\n");
17     pthread_create(&p1, NULL, mythread, "A");
18     pthread_create(&p2, NULL, mythread, "B");
19     // join waits for the threads to finish
20     pthread_join(p1, NULL);
21     pthread_join(p2, NULL);
22     printf("main: end\n");
23     return 0;
24 }

```

```
#include <pthread.h>
#include <stdio.h>

void *thread_fn(void *arg){
    long id = (long) arg;
    printf("Starting thread %ld\n", id);
    sleep(5);
    printf("Exiting thread %ld\n", id);
    return NULL;
}

int main(){
    pthread_t t1, t2;

    pthread_create(&t1, NULL, thread_fn, (void *)1);
    pthread_create(&t2, NULL, thread_fn, (void *)2);
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
    printf("Exiting main\n");
    return 0;
}
```

Thread Trace

Single CPU

main	Thread 1	Thread2
starts running prints "main: begin" creates Thread 1 creates Thread 2 waits for T1		
	runs prints "A" returns	
waits for T2		runs prints "B" returns
prints "main: end"		

main	Thread 1	Thread2
starts running prints "main: begin" creates Thread 1		
	runs prints "A" returns	
		creates Thread 2
waits for T1 <i>returns immediately; T1 is done</i> waits for T2 <i>returns immediately; T2 is done</i> prints "main: end"		runs prints "B" returns

main	Thread 1	Thread2
starts running prints "main: begin" creates Thread 1 creates Thread 2		
		runs prints "B" returns
	waits for T1	
		runs prints "A" returns
waits for T2 <i>returns immediately; T2 is done</i> prints "main: end"		