

Example

```
1 #include <stdio.h>
2 #include <assert.h>
3 #include <pthread.h>
4 #include "common.h"
5 #include "common_threads.h"
6
7 void *mythread(void *arg) {
8     printf("%s\n", (char *) arg);
9     return NULL;
10 }
11
12 int
13 main(int argc, char *argv[]) {
14     pthread_t p1, p2;
15     int rc;
16     printf("main: begin\n");
17     Pthread_create(&p1, NULL, mythread, "A");
18     Pthread_create(&p2, NULL, mythread, "B");
19     // join waits for the threads to finish
20     Pthread_join(p1, NULL);
21     Pthread_join(p2, NULL);
22     printf("main: end\n");
23     return 0;
24 }
```

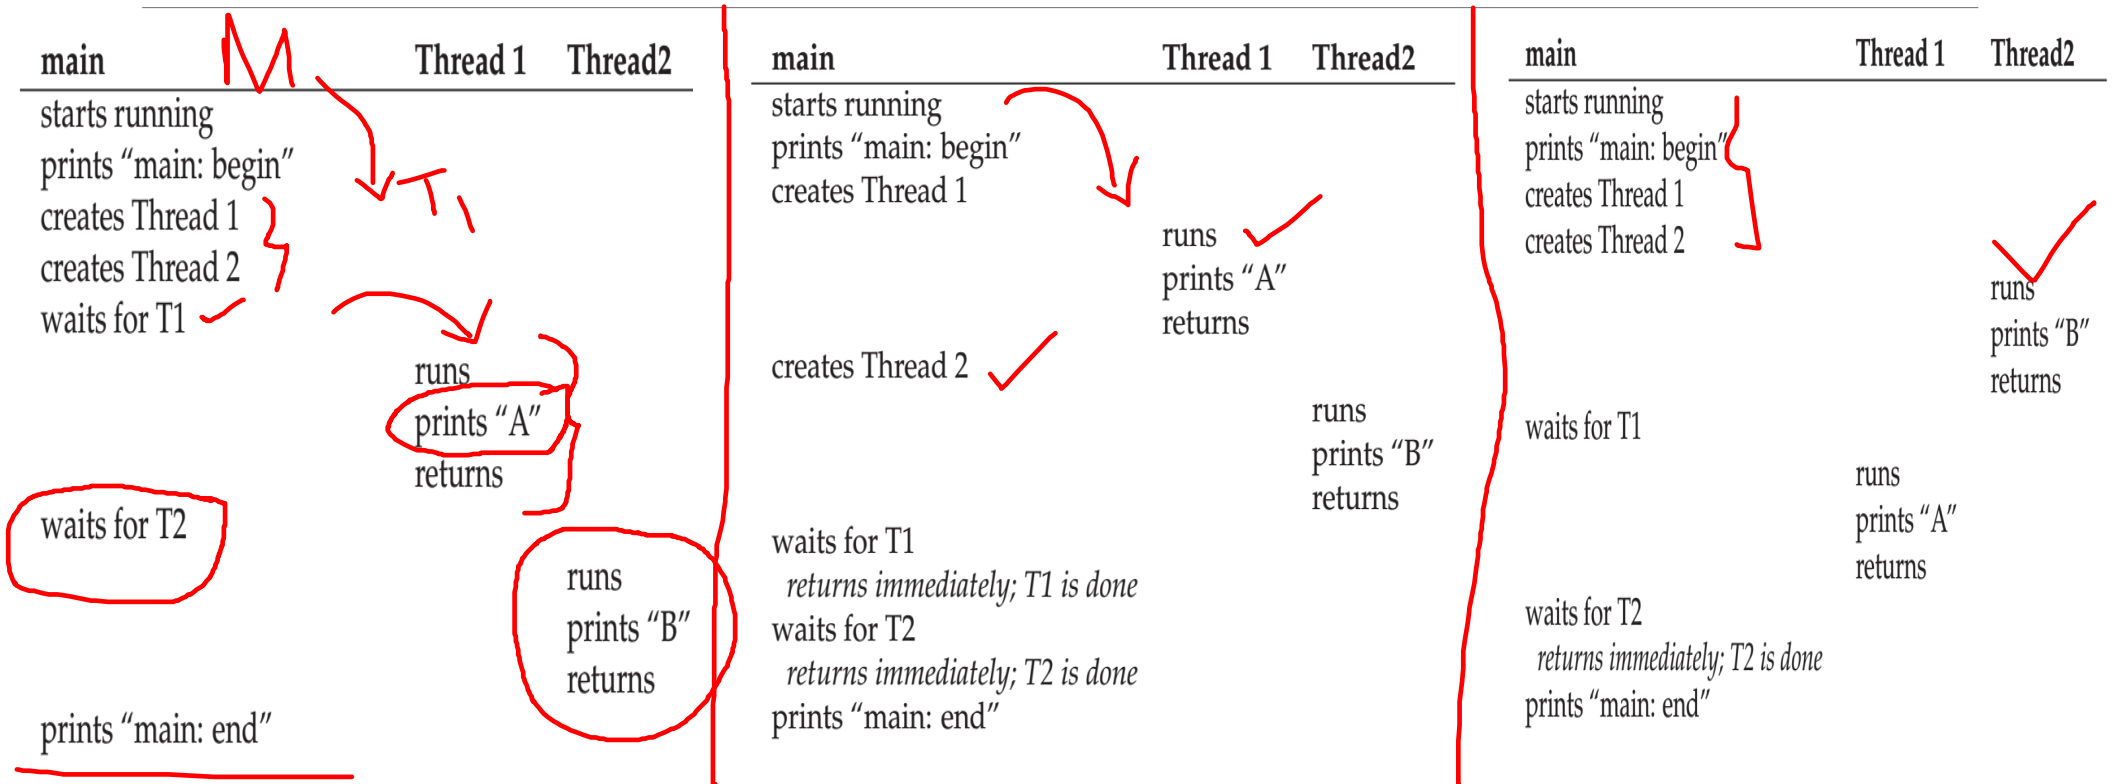
```
#include <pthread.h>
#include <stdio.h>

void *thread_fn(void *arg){
    long id = (long) arg;
    printf("Starting thread %ld\n", id);
    sleep(5);
    printf("Exiting thread %ld\n", id);
    return NULL;
}

int main(){
    pthread_t t1, t2;

    pthread_create(&t1, NULL, thread_fn, (void *)1);
    pthread_create(&t2, NULL, thread_fn, (void *)2);
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
    printf("Exiting main\n");
    return 0;
}
```

Thread Trace



Example: With Shared data

```
1 #include <stdio.h>
2 #include <pthread.h>
3 #include "common.h"
4 #include "common_threads.h"
5
6 static volatile int counter = 0;
7
8 // mythread()
9 //
10 // Simply adds 1 to counter repeatedly, in a loop
11 // No, this is not how you would add 10,000,000 to
12 // a counter, but it shows the problem nicely.
13 //
14 void *mythread(void *arg) {
15     printf("%s: begin\n", (char *) arg);
16     int i;
17     for (i = 0; i < 1e7; i++) {
18         counter = counter + 1;
19     }
20     printf("%s: done\n", (char *) arg);
21     return NULL;
22 }
23
24 // main()
25 //
26 // Just launches two threads (pthread_create)
27 // and then waits for them (pthread_join)
28 //
29 int main(int argc, char *argv[]) {
30     pthread_t p1, p2;
31     printf("main: begin (counter = %d)\n", counter);
32     pthread_create(&p1, NULL, mythread, "A");
33     pthread_create(&p2, NULL, mythread, "B");
34
35     // join waits for the threads to finish
36     pthread_join(p1, NULL);
37     pthread_join(p2, NULL);
38     printf("main: done with both (counter = %d)\n",
39           counter);
40     return 0;
41 }
```

Mutual
Exclusion

Mutex

c=0
c=?

-lpthread

With shared data – What happens?

```
prompt> gcc -o main main.c -Wall -pthread; ./main
main: begin (counter = 0)
A: begin ✓
B: begin ✓
A: done ✓
B: done ✓
main: done with both (counter = 20000000)
```

```
prompt> ./main
main: begin (counter = 0)
A: begin
B: begin
A: done
B: done
main: done with both (counter = 19345221)
```

What do we expect? Two threads, each increments counter by 10^7 , so 2×10^7

Sometimes, a lower value. Why?

*Objdump -d main
disassembler*

Assembly code

OS	Thread 1	Thread 2	(after instruction)		
			PC	eax	counter
	<i>before critical section</i>		100	0	50
	mov 8049a1c, %eax		105	50	50
	add \$0x1, %eax		108	51	50
interrupt					
save T1					
restore T2					
		mov 8049a1c, %eax	100	0	50
		add \$0x1, %eax	105	50	50
		mov %eax, 8049a1c	108	51	50
			113	51	51
interrupt					
save T2					
restore T1					
	mov %eax, 8049a1c		108	51	51
			113	51	51

Assembly code of

counter = counter + 1

100 mov 0x8049a1c, %eax
 105 add \$0x1, %eax
 108 mov %eax, 0x8049a1c

52

Race condition & Synchronization

What just happened is called a race condition –Concurrent execution can lead to different results


- Critical section: portion of code that can lead to race conditions
- What we need: mutual exclusion–Only one thread should be executing critical section at any time
- What we need: atomicity of the critical section –The critical section should execute like one uninterruptible instruction
- How is it achieved? Locks

Who manages threads?

- Two strategies
 - User threads
 - Thread management done by user level thread library. Kernel knows nothing about the threads.
 - Kernel threads
 - Threads directly supported by the kernel.
 - Known as light weight processes.


User level threads

- **Advantages:**
 - Fast (really lightweight)
(no system call to manage threads. The thread library does everything).
 - Can be implemented on an OS that does not support threading.
 - Switching is fast. No, switch from user to protected mode.
- **Disadvantages:**
 - Scheduling can be an issue. (Consider, one thread that is blocked on an IO and another runnable.)
 - Lack of coordination between kernel and threads. (A process with 1000 threads competes for a timeslice with a process having just 1 thread.)
 - Requires non-blocking system calls. (If one thread invokes a system call, all threads need to wait)




Kernel level threads

- **Advantages:**
 - Scheduler can decide to give more time to a process having large number of threads than process having small number of threads.
 - Kernel-level threads are especially good for applications that frequently block.
- **Disadvantages:**
 - The kernel-level threads are slow (they involve kernel invocations.)
 - Overheads in the kernel. (Since kernel must manage and schedule threads as well as processes. It requires a full thread control block (TCB) for each thread to maintain information about threads.)




Many-to-One-Model

- Many user level threads map to a single kernel thread
- Pros:
 - Fast. No system calls to manage threads.
 - No mode change for switching threads
- Cons:
 - No parallel execution of threads. All threads block when one has a system call.
 - Not suited for multi-processor systems.




One-to-One Model

- Each user thread associated with one kernel thread.
- **Pros.**
 - Better suited for multiprocessor environments.
 - When one thread blocks, the other threads can continue to execute.
- **Cons.**
 - Expensive. Kernel is involved.



Many-to-many

- Many user threads mapped to many kernel threads
 - Supported by some unix and windows versions
- **Pros:** flexible
 - OS creates kernel threads as required
 - Process creates user threads as needed
- **Cons:** Complex
 - Double management



Summary

- A **critical section** is a piece of code that accesses a *shared* resource, usually a variable or data structure.
- A **race condition** (or **data race** [NM92]) arises if multiple threads of execution enter the critical section at roughly the same time; both attempt to update the shared data structure, leading to a surprising (and perhaps undesirable) outcome.
- An **indeterminate** program consists of one or more race conditions; the output of the program varies from run to run, depending on which threads ran when. The outcome is thus not **deterministic**, something we usually expect from computer systems.
- To avoid these problems, threads should use some kind of **mutual exclusion** primitives; doing so guarantees that only a single thread ever enters a critical section, thus avoiding races, and resulting in deterministic program outputs.

Any solution should satisfy the following requirements

- Mutual Exclusion : No more than one process in critical section at a given time
- Progress : When no process is in the critical section, any process that requests entry into the critical section must be permitted without any delay }
- No starvation (bounded wait): There is an upper bound on the number of times a process enters the critical section, while another is waiting

Locks

Consider update of shared variable **balance = balance + 1**

We can use a special lock variable to protect it

```
lock_t mutex; // some globally-allocated lock 'mutex'  
...  
lock(&mutex);  
balance = balance + 1;  
unlock(&mutex);
```

All threads accessing a critical section share a lock

- One thread succeeds in locking – owner of lock ✓
- Other threads that try to lock cannot proceed further until lock is released by the owner
- Pthreads library in Linux provides such locks ✓

Waiting (Blocked)

T_1 ✓
 program 0
 {
 *
 *
 lock(L) ✓
 counter++
 unlock(L) ✓
 *
 }

shared variable
 int counter=5;
 lock_t L;

T_2
 program 1
 {
 *
 *
 lock(L) ✓
 counter--
 unlock(L)
 *
 }


- lock(L) acquire lock L exclusively
 - Only the process with L can access the critical section
- unlock(L) : release exclusive access to lock L
 - Permitting other processes to access the critical section

Building a Lock

Goals of a lock implementation

- Mutual exclusion ✓
- Fairness: all threads should eventually get the lock, and no thread should starve
- Low overhead: acquiring, releasing, and waiting for lock should not consume too many resources
 - Implementation of locks are needed for both userspace programs (e.g., pthreads library) and kernel code
 - Implementing locks needs support from hardware and OS

Using Interrupts



- Simple
 - When interrupts are disabled, context switches won't happen
- Requires privileges
 - User processes generally cannot disable interrupts
- Not suited for multicore systems ✓

Is disabling interrupts enough ?

Is this enough?

- No, not always!
- Many issues here:
 - Disabling interrupts is a privileged instruction and program can misuse it (e.g., run forever)
 - Will not work on multiprocessor systems, since another thread on another core can enter critical section
- This technique is used to implement locks on single processor systems inside the OS
 - Need better solution for other situations

Implementation Attempt -1

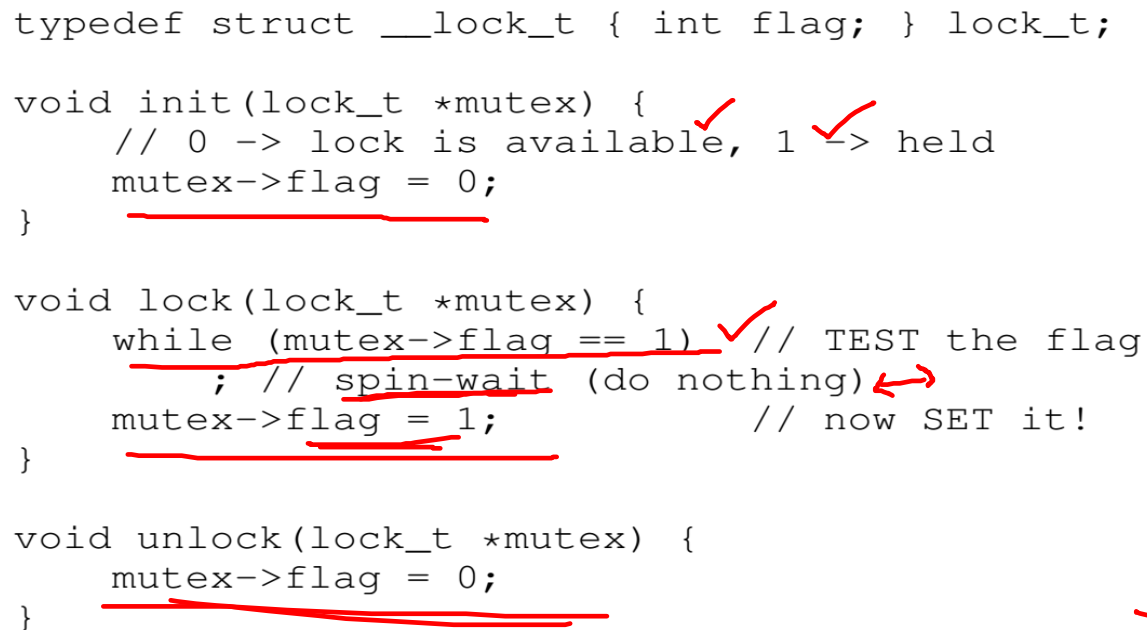
- Lock: spin on a flag variable until it is unset, then set it to acquire lock
- Unlock: unset flag variable

```
typedef struct __lock_t { int flag; } lock_t;

void init(lock_t *mutex) {
    // 0 -> lock is available, 1 -> held
    mutex->flag = 0;
}

void lock(lock_t *mutex) {
    while (mutex->flag == 1) // TEST the flag
    ; // spin-wait (do nothing)
    mutex->flag = 1; // now SET it!
}

void unlock(lock_t *mutex) {
    mutex->flag = 0;
}
```



- Thread 1 spins, lock is released, ends spin
- Thread 1 interrupted just before setting flag
- Race condition has moved to the lock acquisition code!

Thread 1

call lock () ✓

while (flag == 1) ✓

interrupt: switch to Thread 2

flag = 1; // set flag to 1 (too!)

Thread 2

call lock () ✓

while (flag == 1) ✓

flag = 1; ✓

interrupt: switch to Thread 1